

The Auditable Extraction Checklist

10 questions to ask before you trust an automated tool to pull data out of your documents — for legal, finance, claims, healthcare and research teams.

Automating document data entry is easy to buy and easy to regret. The tools are fast and mostly right — and "mostly right, with no way to check" is exactly how a wrong number reaches your ledger, your case file, or your patient record. Use this checklist to separate a tool you can defend from a black box you'll have to apologise for.

1 Can you see the source of every extracted value?

☐

For each field the tool returns, can you trace it back to the exact location in the original document — not just "somewhere in the file", but the specific line, cell or sentence?

2 How long does it take to verify one field?

☐

If checking a value means re-reading the document, people won't do it. The bar to clear: verification is a single click, not a re-read.

3 Does it ever invent values that aren't in the document?

☐

Chat-based tools can produce plausible-but-fabricated data. Confirm the tool only returns text that actually appears in the source — and that you can prove it.

4 Is the extraction repeatable?

☐

Run the same document twice — do you get the same fields? Ad-hoc prompts drift. Saved templates or schemas keep your team extracting the same way every time.

5 Can you export the data *with* its citations?

☐

A spreadsheet of unexplained values isn't auditable. You want structured output (JSON/CSV) where each value carries a pointer back to its source.

6 What happens to your documents after processing?

☐

Are they used to train models? How long are they retained? Is there a Data Processing Agreement? Get the answers in writing before you upload anything sensitive.

7 Who can catch an error — and when?

☐

The right time to catch a bad value is before it's committed, by the person responsible for it. A tool that only surfaces errors in a downstream audit has moved the risk, not removed it.

- 8

Does it handle your real documents, not just clean ones?

☐
- Test it on your messiest invoice, your densest contract, your worst scan. Accuracy on a demo file tells you nothing about Tuesday.
- 9

Can it scale from one document to a batch?

☐
- One-off extraction is a party trick. Check for batch processing and an API so the tool fits an actual workflow, not just a manual copy-paste replacement.
- 10

If a regulator or client asked "where did this come from?", could you answer?

☐
- This is the whole test. If the answer is "the AI said so", the tool has failed. If the answer is "here's the exact span in the source", you have an auditable process.

Why "source grounding" is the line that matters

Every question above collapses into one idea: an answer you can trace is an answer you can trust. Source grounding — anchoring each extracted field to its exact span in the original document — is what turns automated extraction from a black box into a defensible process. It's the difference between speed and reckless speed.

Groundtruth grounds every field by default.

Extract structured data from any document and verify each value in one click. Free plan — 25 pages/month, no card.

[Try it free →](#)